

Contents

1	Basic Test Results	2
2	README	3
3	Barrier.h	4
4	Barrier.cpp	5
5	Makefile	6
6	MapReduceFramework.cpp	7

1 Basic Test Results

```
1  ===== Tar Content Test =====
2  found README
3  found Makefile
4  tar content test PASSED!
5
6  ===== logins =====
7  login names mentioned in file:  kawabanga,michal.maayan
8  Please make sure that these are the correct login names.
9
10 ===== make Command Test =====
11 g++ -Wall -std=c++11 -g -I. -c -o MapReduceFramework.o MapReduceFramework.cpp
12 g++ -Wall -std=c++11 -g -I. -c -o Barrier.o Barrier.cpp
13 ar rv libMapReduceFramework.a MapReduceFramework.o Barrier.o Barrier.h
14 a - MapReduceFramework.o
15 a - Barrier.o
16 a - Barrier.h
17 ranlib libMapReduceFramework.a
18
19 MapReduceFramework.cpp: In function void* threadLogic(void*):
20 MapReduceFramework.cpp:111:21: warning: comparison between signed and unsigned integer expressions [-Wsign-compare]
21     while (oldValue < tc->inputVec->size()) {
22         ~~~~~
23 MapReduceFramework.cpp:131:41: warning: comparison between signed and unsigned integer expressions [-Wsign-compare]
24     for (int i = FIRSTTHREAD; i < tc->MT; ++i) {
25         ~~~~~
26 MapReduceFramework.cpp:159:37: warning: comparison between signed and unsigned integer expressions [-Wsign-compare]
27     for (int i = FIRSTTHREAD; i < tc->MT; ++i) {
28         ~~~~~
29 MapReduceFramework.cpp:172:19: warning: comparison between signed and unsigned integer expressions [-Wsign-compare]
30     if (index < tc->Queue->size()) {
31         ~~~~~
32 MapReduceFramework.cpp: In function void runMapReduceFramework(const MapReduceClient&, const InputVec&, OutputVec&, int):
33 MapReduceFramework.cpp:218:55: warning: narrowing conversion of i from int to unsigned int inside { } [-Wnarrowing]
34         &Queue, &mutexQueue, &fillCount};
35         ^
36 ar: creating libMapReduceFramework.a
37
38 stderr: b'MapReduceFramework.cpp: In function void* threadLogic(void*)\nMapReduceFramework.cpp:111:
39 make command test FAILED!
40
41 Pre-submission Test FAILED!
42 Check info above.
```

2 README

```
1  kawabanga, michal.maayan
2  Yaron Scherf (305758211), Michal Maayan (203348776)
3  EX: 3
4
5  FILES:
6  MapReduceFramework.c -- Our main code file, an implementation of the
7                          required code.
8  Barrier.h -- school implementation
9  Barrier.cpp -- school implementation
10
11 REMARKS:
12 Hope you have enjoyed from our implementation, we have done our best :)
```

3 Barrier.h

```
1  #ifndef BARRIER_H
2  #define BARRIER_H
3  #include <pthread.h>
4
5  // a multiple use barrier
6
7  class Barrier {
8  public:
9      Barrier(int numThreads);
10     ~Barrier();
11     void barrier();
12
13 private:
14     pthread_mutex_t mutex;
15     pthread_cond_t cv;
16     int count;
17     int numThreads;
18 };
19
20 #endif //BARRIER_H
```

4 Barrier.cpp

```
1  #include "Barrier.h"
2  #include <cstdlib>
3  #include <stdio>
4
5  Barrier::Barrier(int numThreads)
6      : mutex(PTHREAD_MUTEX_INITIALIZER)
7      , cv(PTHREAD_COND_INITIALIZER)
8      , count(0)
9      , numThreads(numThreads)
10 { }
11
12
13 Barrier::~Barrier()
14 {
15     if (pthread_mutex_destroy(&mutex) != 0) {
16         fprintf(stderr, "[Barrier] error on pthread_mutex_destroy");
17         exit(1);
18     }
19     if (pthread_cond_destroy(&cv) != 0){
20         fprintf(stderr, "[Barrier] error on pthread_cond_destroy");
21         exit(1);
22     }
23 }
24
25
26 void Barrier::barrier()
27 {
28     if (pthread_mutex_lock(&mutex) != 0){
29         fprintf(stderr, "[Barrier] error on pthread_mutex_lock");
30         exit(1);
31     }
32     if (++count < numThreads) {
33         if (pthread_cond_wait(&cv, &mutex) != 0){
34             fprintf(stderr, "[Barrier] error on pthread_cond_wait");
35             exit(1);
36         }
37     } else {
38         count = 0;
39         if (pthread_cond_broadcast(&cv) != 0) {
40             fprintf(stderr, "[Barrier] error on pthread_cond_broadcast");
41             exit(1);
42         }
43     }
44     if (pthread_mutex_unlock(&mutex) != 0) {
45         fprintf(stderr, "[Barrier] error on pthread_mutex_unlock");
46         exit(1);
47     }
48 }
```

5 Makefile

```
1  CC=g++
2  CXX=g++
3  RANLIB=ranlib
4
5  LIBSRC=MapReduceFramework.cpp Barrier.cpp Barrier.h
6  LIBOBJ=$(LIBSRC:.cpp=.o)
7
8  INCS=-I.
9  CFLAGS = -Wall -std=c++11 -g $(INCS)
10 CXXFLAGS = -Wall -std=c++11 -g $(INCS)
11
12 OSMLIB = libMapReduceFramework.a
13 TARGETS = $(OSMLIB)
14
15 TAR=tar
16 TARFLAGS=-cvf
17 TARNAME=ex3.tar
18 TARSRC=$(LIBSRC) Makefile README
19
20 all: $(TARGETS)
21
22 $(TARGETS): $(LIBOBJ)
23     $(AR) $(ARFLAGS) $@ $^
24     $(RANLIB) $@
25
26 clean:
27     $(RM) $(TARGETS) $(OSMLIB) $(OBJ) $(LIBOBJ) *~ *core
28
29 depend:
30     makedepend -- $(CFLAGS) -- $(SRC) $(LIBSRC)
31
32 tar:
33     $(TAR) $(TARFLAGS) $(TARNAME) $(TARSRC)
```

6 MapReduceFramework.cpp

```
1  #include <cstdlib>
2  #include <cstdio>
3  #include "Barrier.h"
4  #include "MapReduceClient.h"
5  #include <atomic>
6  #include <iostream>
7  #include <algorithm>    // std::sort
8  #include <semaphore.h>
9
10 #define FIRSTTHREAD 0
11
12 typedef struct ThreadContext{
13     unsigned int threadId;
14     unsigned int MT;
15     Barrier * barrier;
16     std::atomic<int>* atomicIndex;
17     std::atomic<int> *outAtomicIndex;
18     std::atomic<int> *reduceAtomic;
19     const InputVec* inputVec;
20     OutputVec* outputVec;
21     std::vector<IntermediateVec> *arrayOfInterVec;
22     const MapReduceClient* client;
23     std::vector<IntermediateVec> *Queue;
24     sem_t *mutexQueue;
25     sem_t *fillCount;
26 }ThreadContext;
27
28 typedef struct MapContext{
29     IntermediateVec *interVector;
30 }MapContext;
31
32 typedef struct ReduceContext{
33     OutputVec *outVector;
34     std::atomic<int> *outAtomicIndex;
35 }ReduceContext;
36
37 void safeExit(ThreadContext* tc)
38 {
39     if(tc == nullptr){
40         exit(-1);
41     }
42     tc->Queue->clear();
43 }
44
45 void printErr(std::string msg, ThreadContext *tc){
46     std::cerr<<msg<<std::endl;
47     safeExit(tc);
48 }
49
50 //for debug
51 //void printInterVector(IntermediateVec** array, int numOfThreads){
52 //    printf("+++++++\n");
53 //    for (int i = 1; i < numOfThreads; ++i){
54 //        printf("thread id: %d:\n",i);
55 //        for (unsigned long j = 0; j < (*(array[i])).size() ; ++j){
56 //            //printIntermediatePair((*(array[i])).at(j));
57 //        }
58 //        printf("~~~~~\n");
59 //    }
```

```

60  //}
61
62  void emit2 (K2* key, V2* value, void* context){
63      auto* tc = (MapContext*) context;
64      tc->interVec->emplace_back(key, value);
65  }
66  void emit3 (K3* key, V3* value, void* context){
67      auto* tc = (ReduceContext*) context;
68      int oldValue = (*(tc->outAtomicIndex))++ ;
69      auto iterator = (*(tc->outVector)).begin();
70      (*(tc->outVector)).emplace(iterator+(oldValue), key, value);
71  }
72
73  bool comparator(IntermediatePair &p1, IntermediatePair &p2){
74      return (*p1.first < *p2.first);
75  }
76
77  // check the IntermediateVec isn't empty, in case it doesn't find the next max key
78  bool check_empty_find_max(std::vector<IntermediateVec> *arr, int MT, K2 **max){
79      bool isEmpty = true;
80      for (int i = FIRSTTHREAD; i < MT; ++i) {
81          if (not((arr->at(i)).empty())) {
82              isEmpty = false;
83              if(*max == nullptr){
84                  *max = ((arr->at(i)).back().first;
85              }
86              else{
87                  if((**max) < *((arr->at(i)).back().first)
88                  {
89                      *max = ((arr->at(i)).back().first;
90                  }
91              }
92          }
93      }
94      return isEmpty;
95  }
96
97  bool is_eq(K2 *max, IntermediatePair &p){
98      if(not(*max < *p.first)){
99          if(not(*p.first < *max)){
100              return true;
101          }
102      }
103      return false;
104  }
105
106  void* threadLogic (void* context) {
107      auto *tc = (ThreadContext *) context;
108      int oldValue = (*(tc->atomicIndex))++;
109
110      //map logic
111      while (oldValue < tc->inputVec->size()) {
112          auto k1 = tc->inputVec->at(oldValue).first;
113          auto v1 = tc->inputVec->at(oldValue).second;
114          MapContext mapContext = {(&tc->arrayOfInterVec->at(tc->threadId))};
115          tc->client->map(k1, v1, &mapContext);
116          oldValue = (*(tc->atomicIndex))++;
117      }
118
119      //sort
120      auto tempVec = &(tc->arrayOfInterVec->at(tc->threadId)); //[tc->threadId];
121      std::sort(tempVec->begin(), tempVec->end(), comparator);
122      tc->barrier->barrier();
123
124      //shuffle
125      if(tc->threadId == FIRSTTHREAD) {
126          // initial K2
127          K2* max = nullptr;

```



```

128     bool isEmpty = check_empty_find_max(tc->arrayOfInterVec, tc->MT, &max);
129     while(not isEmpty){
130         IntermediateVec sameKey = {};
131         for (int i = FIRSTTHREAD; i < tc->MT; ++i) {
132             // in case the vector isn't empty
133             if (not(tc->arrayOfInterVec->at(i).empty())) {
134                 while(is_eq(max, (tc->arrayOfInterVec->at(i).back())) {
135                     (sameKey).emplace_back(((tc->arrayOfInterVec->at(i)).back()));
136                     ((tc->arrayOfInterVec->at(i))).pop_back();
137                     //check emptyness again
138                     if (((tc->arrayOfInterVec->at(i))).empty()){
139                         break;
140                     }
141                 }
142             }
143         }
144         max = nullptr;
145         isEmpty = check_empty_find_max(tc->arrayOfInterVec, tc->MT, &max);
146         //add to queue using semaphore
147         if (sem_wait(tc->mutexQueue) != 0){
148             printErr("sem_wait err",tc);
149         }
150         tc->Queue->push_back(sameKey);
151         if (sem_post(tc->mutexQueue) != 0){
152             printErr("sem_post err",tc);
153         }
154         if (sem_post(tc->fillCount) != 0){
155             printErr("sem_post err",tc);
156         }
157     }
158     //wakeup all the threads who went down
159     for (int i = FIRSTTHREAD; i < tc->MT; ++i) {
160         if (sem_post(tc->fillCount) != 0){
161             printErr("sem_post err",tc);
162         }
163     }
164 } // end of shuffle
165
166 //reduce
167 while(true){
168     if (sem_wait(tc->fillCount) != 0){
169         printErr("sem_wait err",tc);
170     }
171     auto index = (*(tc->reduceAtomic))++ ;
172     if (index < tc->Queue->size()) {
173         if (sem_wait(tc->mutexQueue) != 0){
174             printErr("sem_wait err",tc);
175         }
176         auto pairs = &(tc->Queue->at(index));
177         ReduceContext reduceContext = {tc->outputVec, tc->outAtomicIndex};
178         tc->client->reduce(pairs, &reduceContext);
179         if (sem_post(tc->mutexQueue) != 0){
180             printErr("sem_post err",tc);
181         }
182     } else{
183         break;
184     }
185 }
186 return nullptr;
187 }
188
189
190 void runMapReduceFramework(const MapReduceClient& client,
191                             const InputVec& inputVec, OutputVec& outputVec,
192                             int multiThreadLevel){
193     if(multiThreadLevel < 1){
194         printErr("multiThreadLevel isn't legal", nullptr);
195     }

```

```

196 pthread_t threads[multiThreadLevel];
197 ThreadContext contexts[multiThreadLevel];
198 Barrier barrier(multiThreadLevel);
199 std::atomic<int> atomicIndex(0);
200 std::atomic<int> outAtomicIndex(0);
201 std::atomic<int> reducetAtomic(0);
202 std::vector<IntermediateVec> arrayOfInterVec = {};
203 std::vector<IntermediateVec> Queue;
204 sem_t mutexQueue;
205 sem_t fillCount;
206 if (sem_init(&mutexQueue, 0, 1) != 0){
207     printErr("sem_init failed\n", nullptr);
208 }
209 if (sem_init(&fillCount, 0, 0) != 0){
210     printErr("sem_init failed\n", nullptr);
211 }
212 for (int i = FIRSTTHREAD; i < multiThreadLevel; ++i) {
213     arrayOfInterVec.push_back({});
214 }
215 for (int i = FIRSTTHREAD; i < multiThreadLevel; ++i) {
216     contexts[i] = {i, (unsigned int)multiThreadLevel, &barrier, &atomicIndex, &outAtomicIndex, &reducetAtomic,
217                  &inputVec, &outputVec, &arrayOfInterVec, &client,
218                  &Queue, &mutexQueue, &fillCount};
219 }
220 for (int i = FIRSTTHREAD+1; i < multiThreadLevel; ++i) {
221     if (pthread_create(threads + i, NULL, threadLogic, contexts + i) != 0){
222         printErr("pthread_create failed\n", contexts + i);
223     }
224 }
225 threadLogic(contexts);
226 for (int i = FIRSTTHREAD+1; i < multiThreadLevel; ++i) {
227     if (pthread_join(threads[i], NULL) != 0){
228         printErr("pthread_join failed\n", contexts + i);
229     }
230 }
231
232 }

```